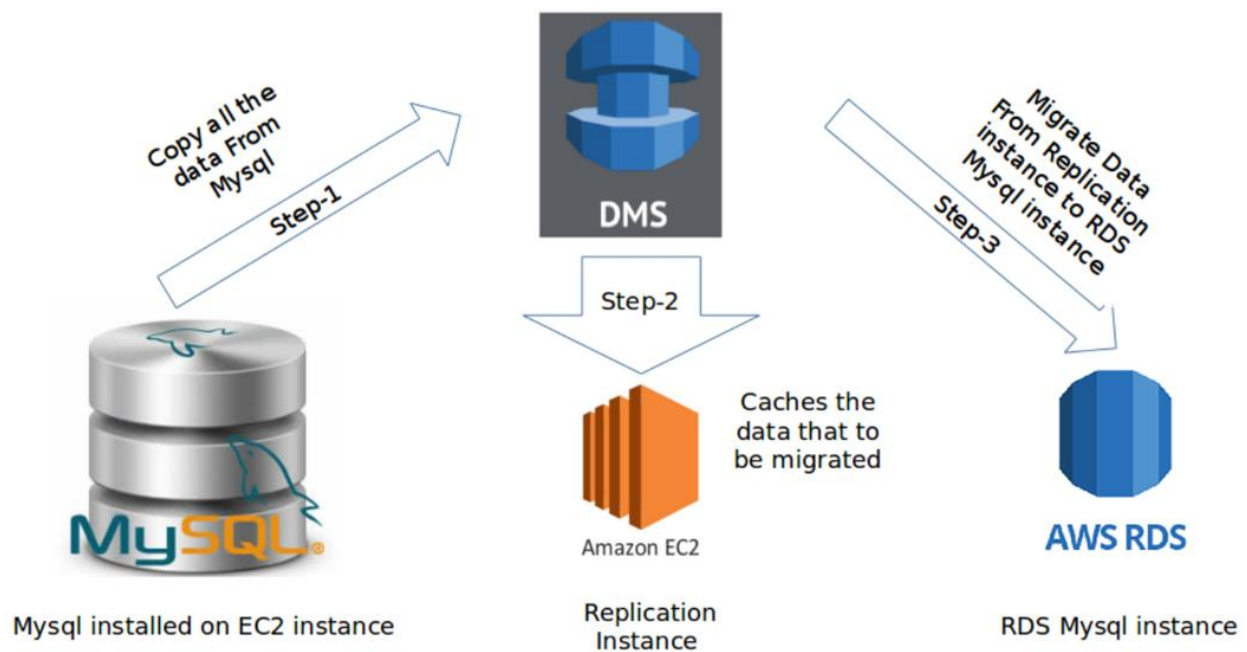


Migrate on-premises MySQL environment to Amazon RDS using Amazon DMS

This project will be executed in 4 different modules:

1. Create a MySQL database instance in Amazon RDS
2. Create a replication instance in AWS DMS
3. Create target and source endpoints in AWS DMS
4. Create a replication task in AWS DMS
5. Complete the migration and clean up the resources

Architectural Diagram



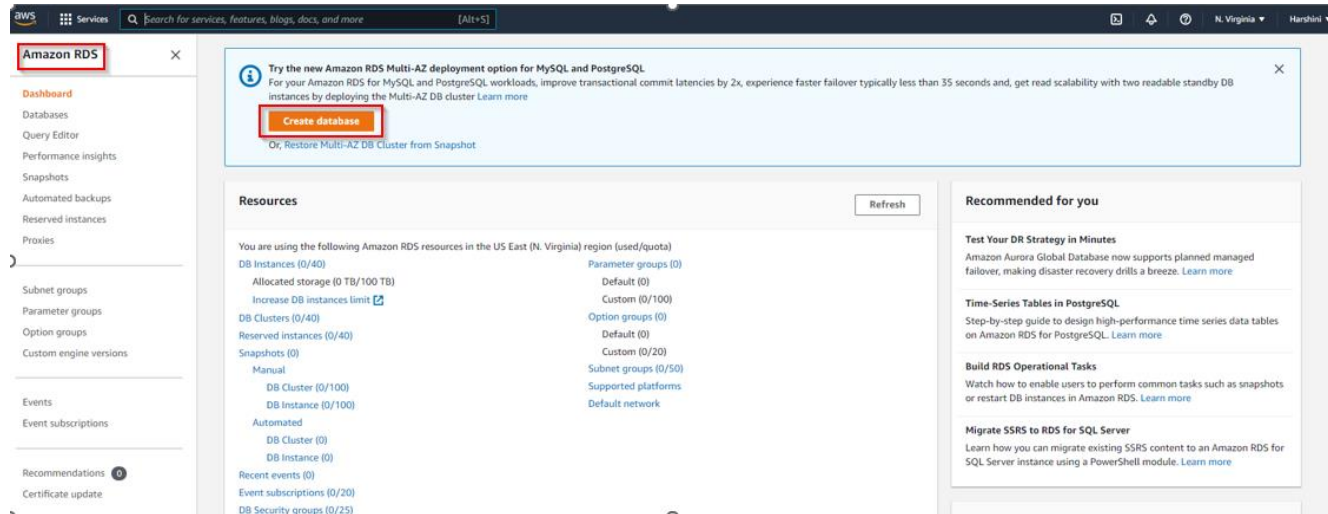
Author: Harshini U
Sr DBA

Module 1: Create a MySQL database instance in Amazon RDS

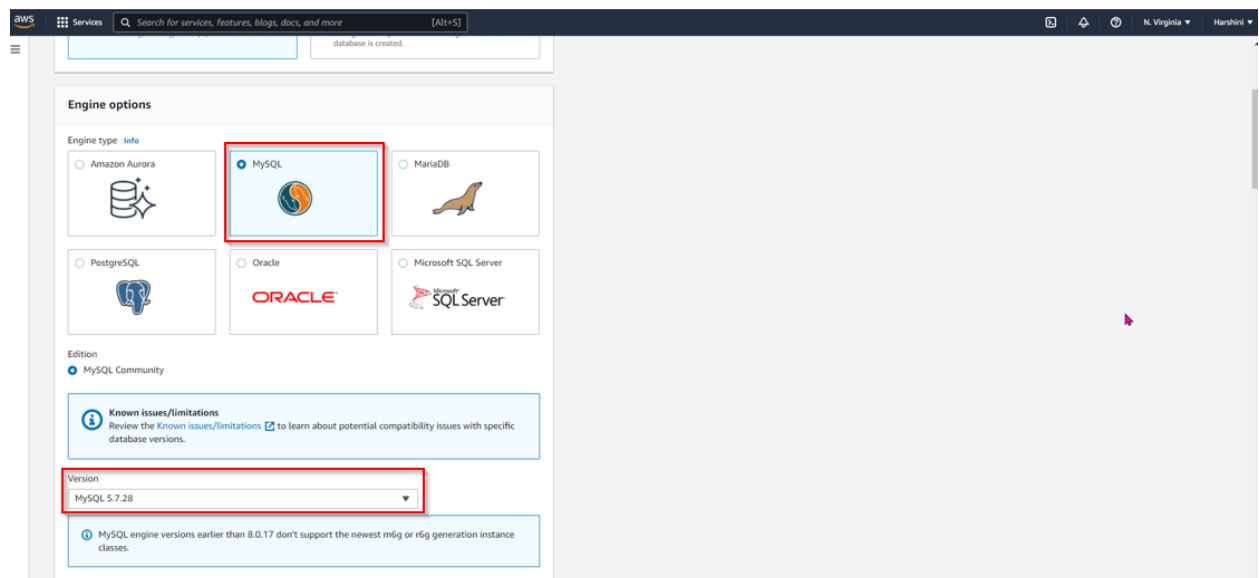
This RDS instance will be used as a primary database after copying existing data into it using AWS DMS.

Step 1: Create an Amazon RDS database instance.

1. Navigate to the **Amazon RDS** console. On the main page, choose **Create database** to create a new database.




2. In the **Engine options** section, choose **MySQL** as the **Engine type**. Then, choose the version **MySQL 5.7.28** to use.



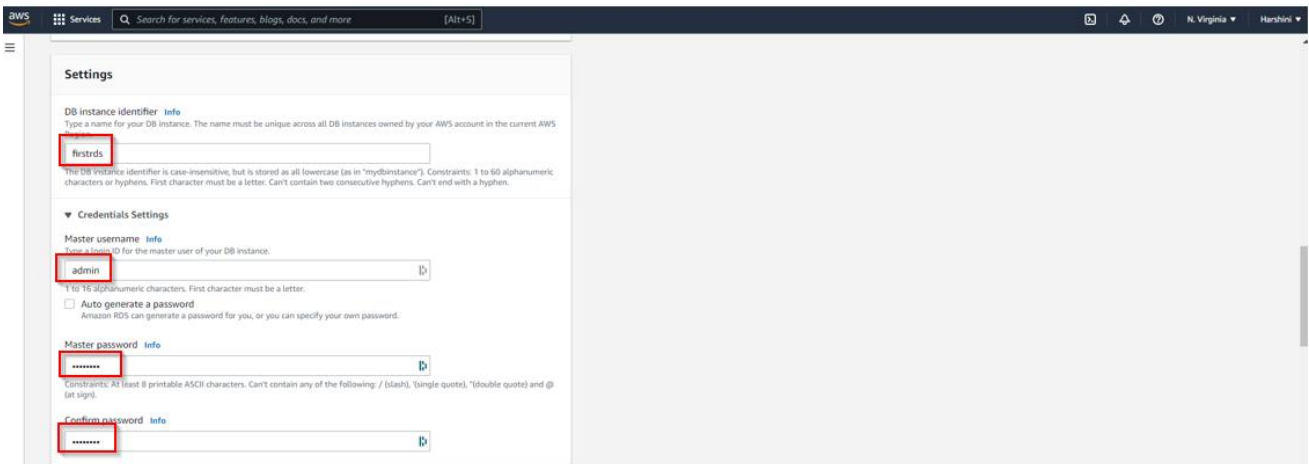
**Author: Harshini U
Sr DBA**

Author: Harshini U
Sr DBA


3. Choose Dev/Test template.



4. In the **Settings** section, database name “**firstrds**” is given and master username “**admin**” and password is set. Saved the master password to connect to the database which can also help to create additional users.

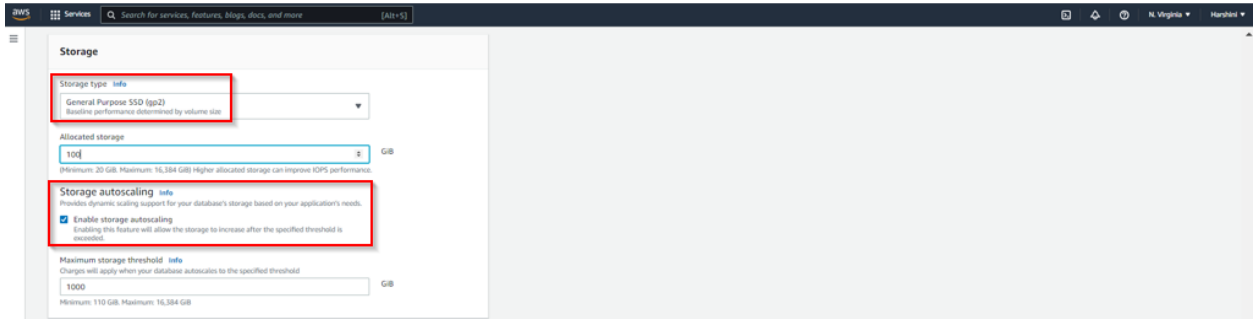


5. Next, database instance size is chosen based on the estimated capacity. Please note that it requires some downtime to change database size.



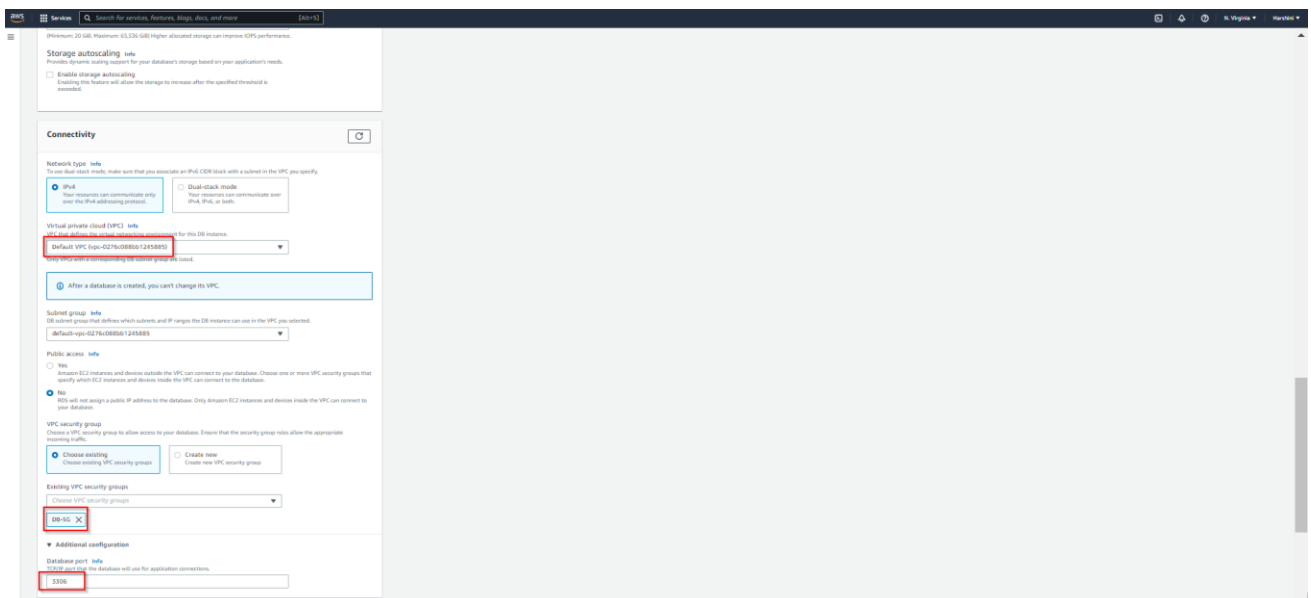
**Author: Harshini U
Sr DBA**

6. Next, configure the storage options for Amazon RDS database.



The screenshot shows the 'Storage' configuration page in the Amazon RDS console. The 'Storage type' is set to 'General Purpose SSD (gp2)'. The 'Allocated storage' is set to '100' GB. The 'Storage autoscaling' section has 'Enable storage autoscaling' checked. The 'Maximum storage threshold' is set to '1000' GB.

7. Configure **Connectivity** - Specify **Amazon Virtual Private Cloud** (Amazon VPC) in which database resides, as well as the network subnet and security groups for database instance.



The screenshot shows the 'Connectivity' configuration page in the Amazon RDS console. The 'Network type' is set to 'IPv4'. The 'Virtual private cloud (VPC)' is set to 'default-vpc-0276c0880a1245885'. The 'Subnet group' is set to 'default-sgs-0276c0880a1245885'. The 'Public access' section has 'No' selected. The 'VPC security group' section has 'Choose existing' selected, and the 'Existing VPC security groups' dropdown is set to 'db-sg-0276c0880a1245885'. The 'Database port' is set to '3306'.

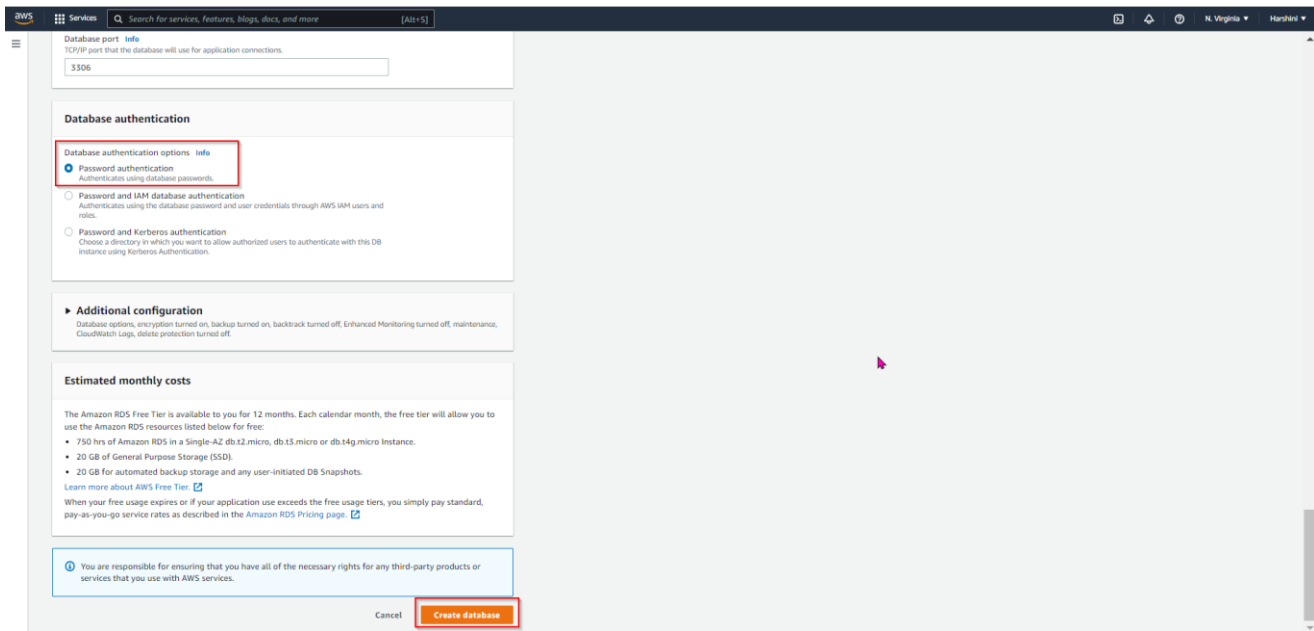
8. Finally, choose the database authentication methods to allow in database as MySQL databases traditionally allow for username and password authentication. With MySQL on Amazon RDS, another option is to allow for authentication by using AWS Identity and Access Management (**IAM**). This integrates easily with your application compute and removes the need for credential rotation.

It is recommended to allow both password and IAM database authentication. It will result in some downtime for Amazon RDS database instance to add IAM database authentication later.


**Author: Harshini U
Sr DBA**

Author: Harshini U
Sr DBA


9. Finally, choose **Create database** to create your database instance.



10. As Amazon RDS is provisioning infrastructure and initializing database, the **Status** of database is **Creating**.



11. When database is ready to use, its **Status** shows **Available** like below.



In this module, successfully created a fully managed MySQL database instance in Amazon RDS.

**Author: Harshini U
Sr DBA**

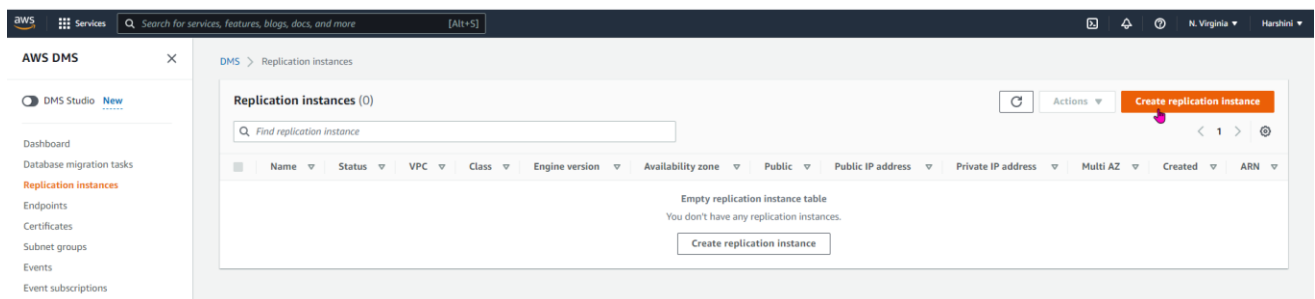
Author: Harshini U
Sr DBA

Module 2: Create a replication instance in AWS DMS

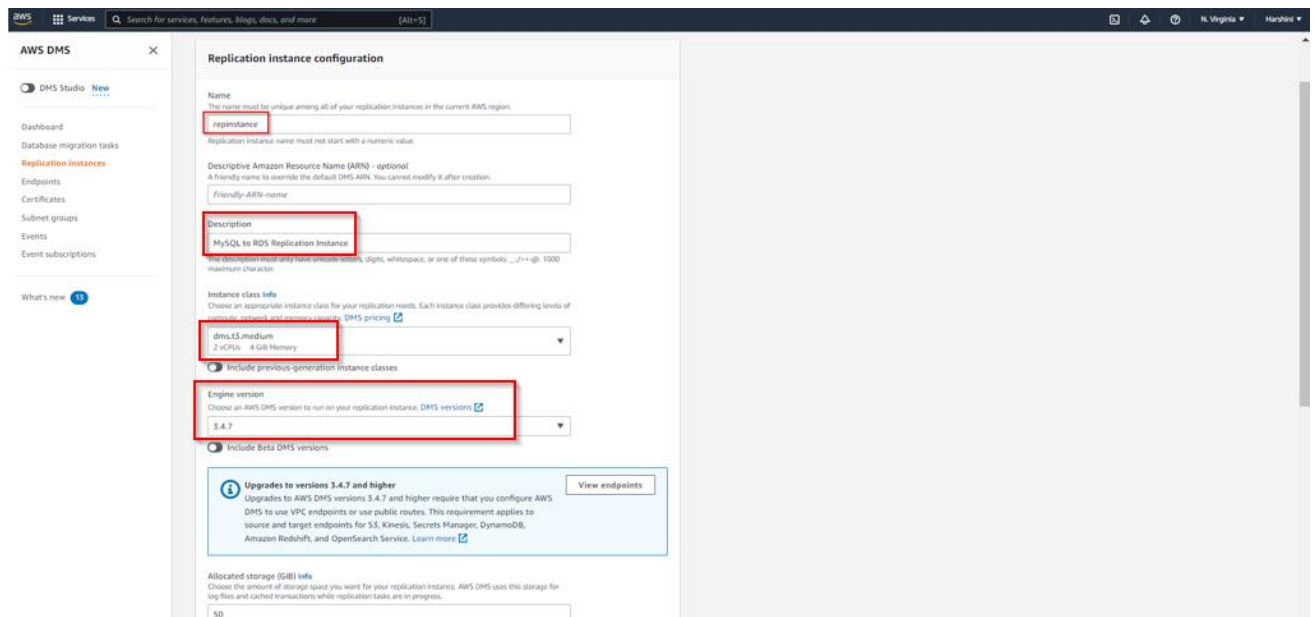
AWS DMS is a service that copies data from an existing database into a fully managed database on AWS. A replication instance is an Amazon EC2 instance that can host replication tasks within AWS DMS.

Step 1: Create a replication instance in AWS DMS.

1. Go to **Replication Instances** section of the AWS DMS console. Choose **Create replication instance** to begin the replication instance creation wizard.



2. In the **Replication instance configuration** section, replication instance name and description is provided. Then instance class is chosen based on the size of an existing database and the amount of data flowing through it. Then, an engine version for AWS DMS is chosen. Finally, the amount of allocated storage for replication instance is also chosen like below.



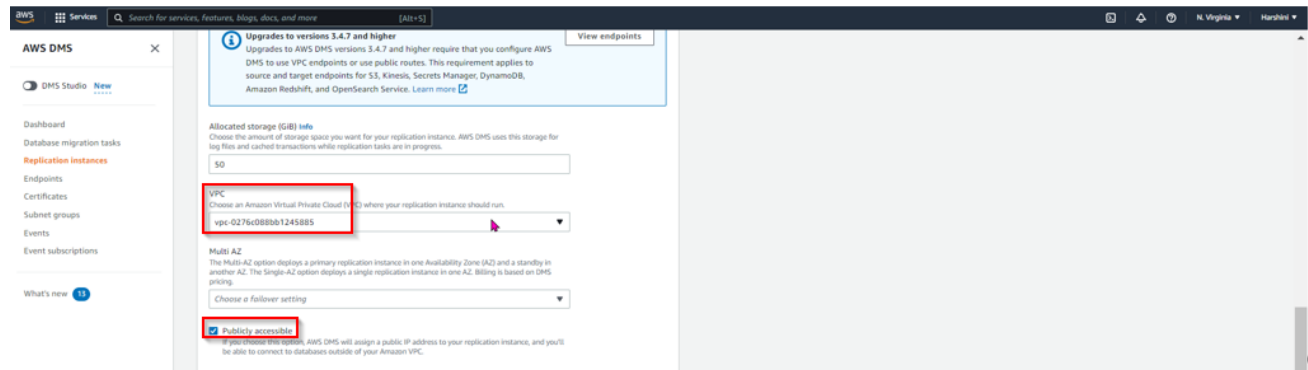
Author: Harshini U
Sr DBA

Author: Harshini U
Sr DBA

3. Choose the same VPC in which you provisioned your Amazon RDS database to ease network access for the replication instance.

Please note that to use a Multi-AZ setup in AWS DMS to keep two databases in sync over a long period of time.

Choose replication instance to be publicly accessible.



4. Next, open **Advanced security and network configuration** section. For the **VPC security group(s)** configuration, choose the same security group that is attached to Amazon RDS database. This allows replication instance to access Amazon RDS database. Choose **Create** to create replication instance in AWS DMS.

**Author: Harshini U
Sr DBA**

Services

Search for services, features, blogs, docs, and more

[Alt+V]

N. Virginia

Harshini

AWS DMS

×

DMS Studio New

Dashboard

Database migration tasks

Replication instances

Endpoints

Certificates

Subnet groups

Events

Event subscriptions

What's new 13

▼ Advanced security and network configuration

Replication subnet group

Choose a subnet group for your replication instance. The subnet group defines the IP ranges and subnets that your replication instance can use within the Amazon VPC you've chosen.

default-vpc-0276c088bb1245885

Availability zone

Choose an availability zone (AZ) where you want your replication instance to run. The default is "No preference", meaning that AWS DMS will determine which AZ to use.

No Preference

VPC security group(s)

Choose one or more security groups for your replication instances. The security group(s) specify inbound and outbound rules to control network access to your replication instance.

Use default

RI-SG X

KMS key info

(Default) aws/dms

Account

Description

Key ARN

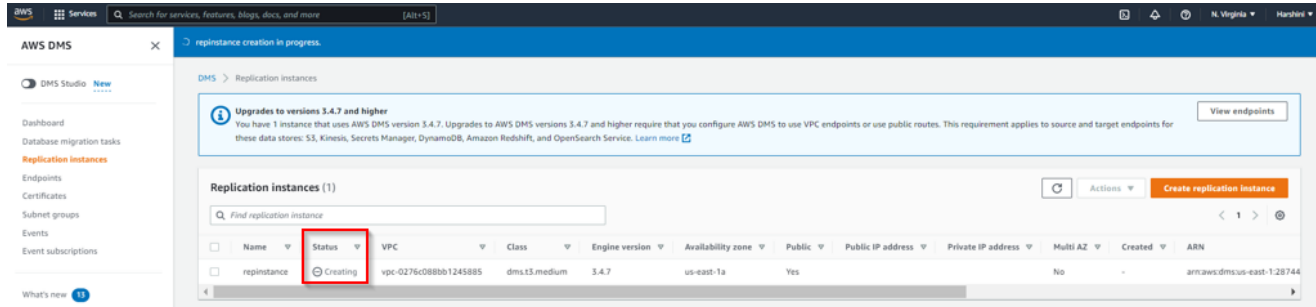
► Maintenance

► Tags

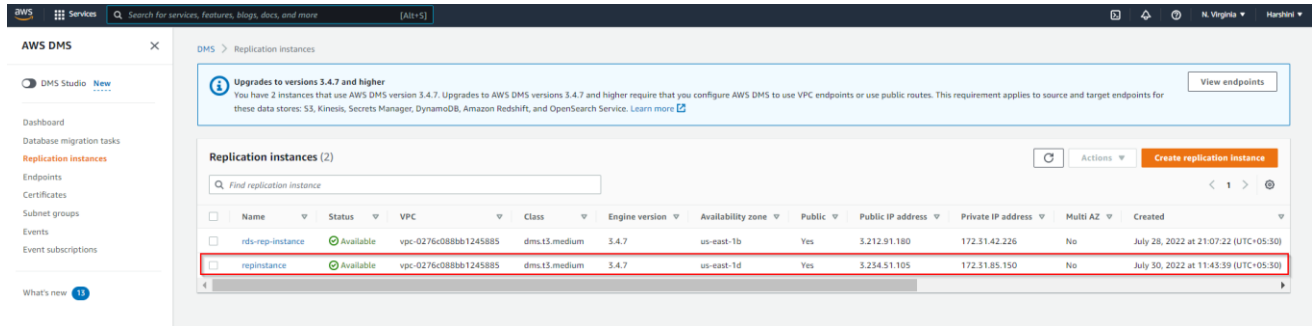
Cancel Create

Author: Harshini U
Sr DBA

5. It shows **Status** of the replication instance as **Creating** while AWS provisions and initializes replication instance.

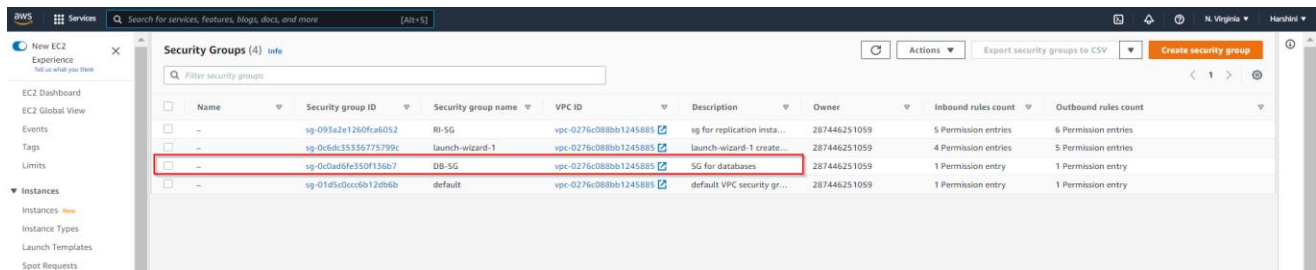


6. When replication instance is ready to go, its **Status** shows **Available**.



Step 2: Add a rule to security group to allow replication instance to access RDS database.

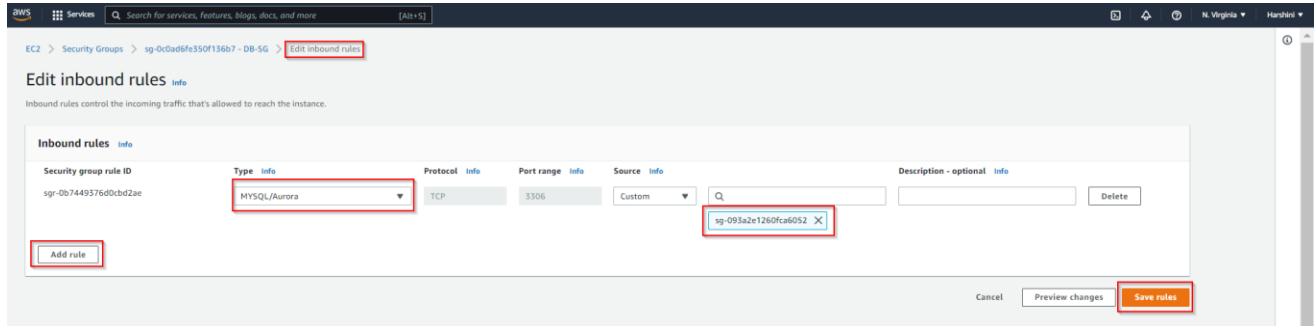
1. In the **Security Groups** section, find the security group which was attached to MySQL database instance, and choose it.



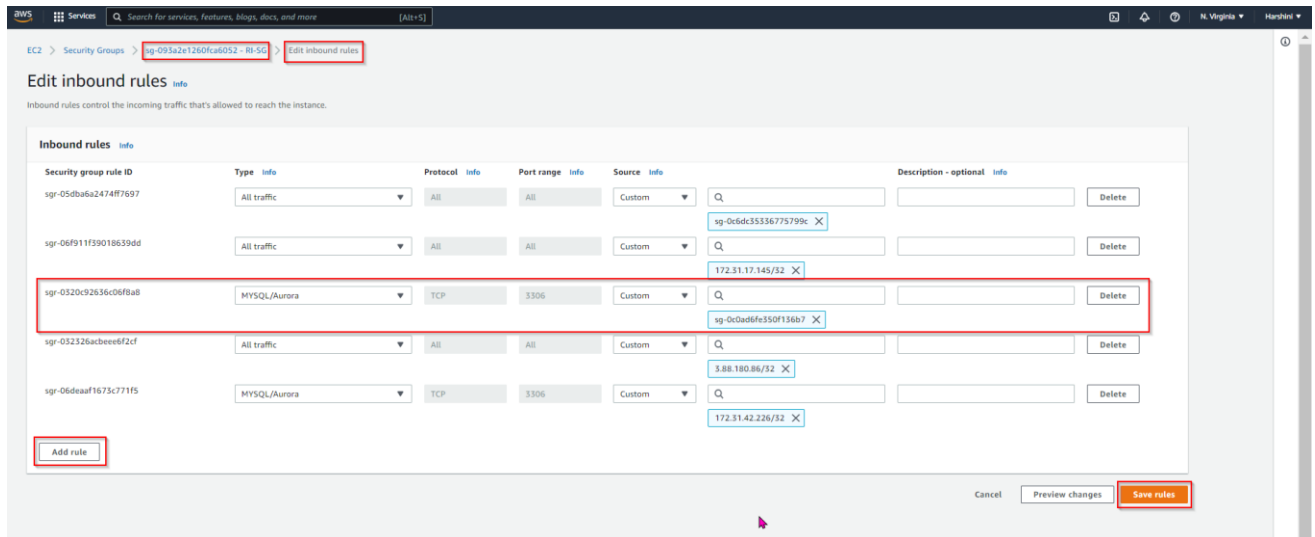
**Author: Harshini U
Sr DBA**

Author: Harshini U
Sr DBA

2. Edit inbound rules of RDS instance security groups and enter the name of the security group used for replication instance like below and choose **Save rules** to save the updated rules for security group.



3. Add database instance security group to replication instance security group by executing same steps as above.



In this module, replication instance is created and also updated the rules for security group.

**Author: Harshini U
Sr DBA**


Module 3: Create endpoints in AWS DMS

A replication task is a job to migrate data from one database to another by using AWS DMS. Before creating a replication task, it is required to register endpoints for source and target databases. An endpoint describes the connection address, credentials, and other information required to connect to a database.

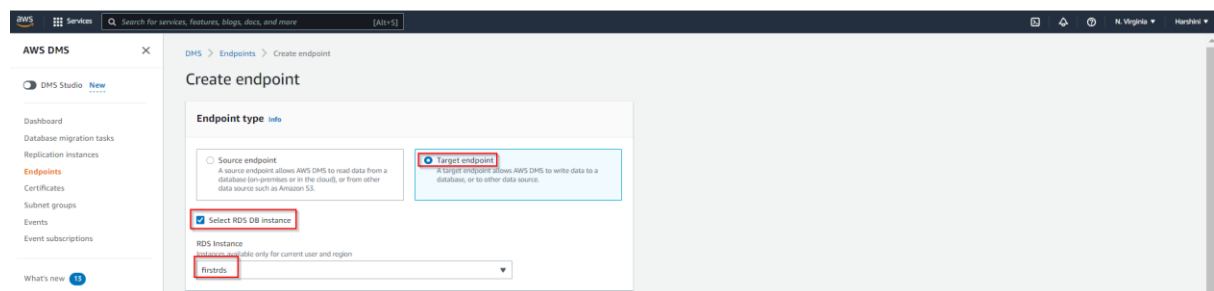
In this module, create source and target endpoints for a replication task in AWS DMS.

Step 1: Create an endpoint for target database (Amazon RDS).

1. Navigate to **Endpoints** section of the AWS DMS console. Choose **Create endpoint** to create a new endpoint.



2. In the endpoint creation wizard, choose to create a **Target endpoint**. Choose the check box to **Select RDS DB instance**, and choose your newly created Amazon RDS database in the dropdown.



**Author: Harshini U
Sr DBA**

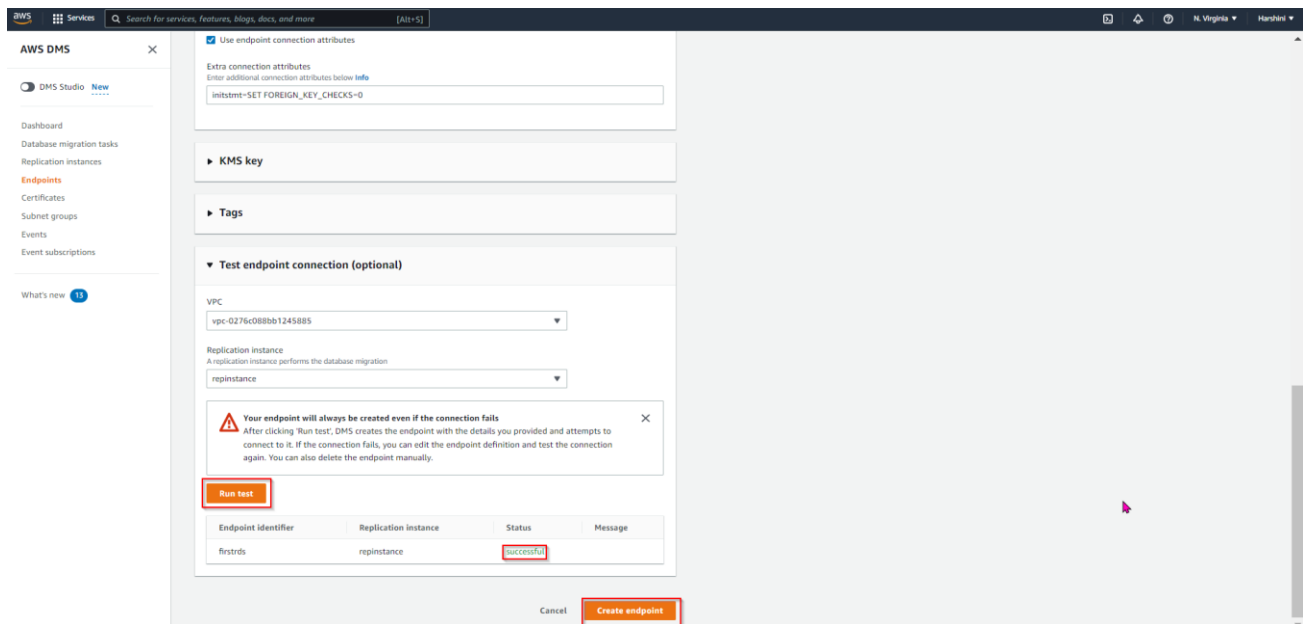
- This fills in most of the **Endpoint configuration** details. It is also required to enter password near the bottom of the section.

- When loading data into a MySQL database by using AWS DMS, it is required to disable foreign key checks, so enter ***initstmt=SET FOREIGN_KEY_CHECKS=0*** in the **Extra connection attributes** box.


- Before saving endpoint, connection is to be tested to ensure that it is configured correctly. To test it, open the **Test endpoint connection** section.

Choose the replication instance in the dropdown list to use, and then select **Run test** and it should show **Status** like **successful**. This indicates that security group and target endpoint are configured correctly.

**Author: Harshini U
Sr DBA**



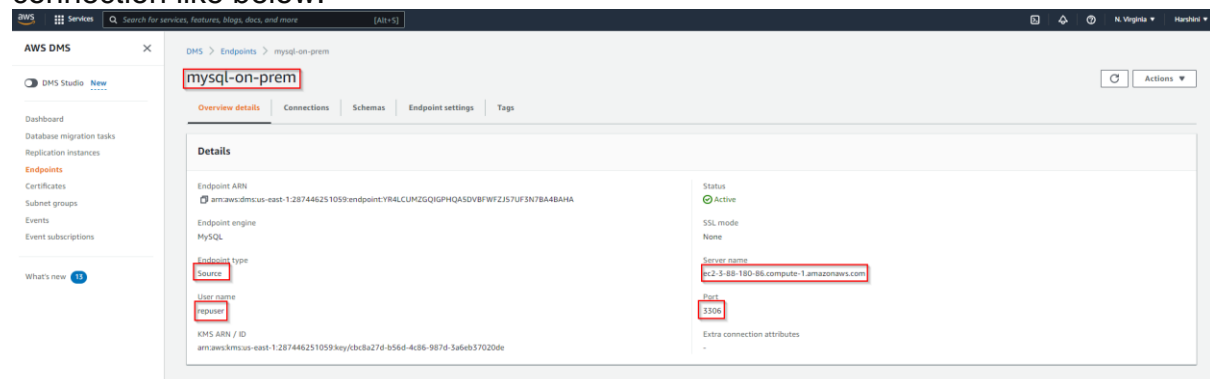
6. To save endpoint, select **Create endpoint**.



Step 2: Configure source endpoint.


It is required to follow the same steps again to create an endpoint for source database. Fill out the connection endpoint, port, and credentials. Provide source database network access to replication instance.

Followed the same steps and created an endpoint for source database and tested connection like below:



**Author: Harshini U
Sr DBA**

Author: Harshini U
Sr DBA



In this module, two endpoints are configured, one for source database and one for target database. Tested both endpoints and successfully able to connect to both databases.


Additional information: For any connectivity issues, please follow this documentation: <https://aws.amazon.com/premiumsupport/knowledge-center/dms-endpoint-connectivity-failures/>

Module 4: Create and run a replication task in AWS DMS

A replication task is responsible for migrating data from a source database to a target database. In this case, I am moving data from an existing MySQL database to newly created database in Amazon RDS.

Step 1: Create a replication task and run it.

1. Navigate to **Database migration tasks** section of the AWS DMS console. Choose **Create task** to create a new replication task.



2. In **Task configuration** section, set up the parameters of the replication task. A task name is given and selected the replication instance created in an earlier module. Then, chose the source endpoint for an existing MySQL database and target endpoint for the fully managed database in Amazon RDS. Chosen **Migrate existing data and replicate ongoing changes** as a migration type with which first type copies all data into RDS database and the second type ensures that all additional updates are replicated to RDS database until the application

**Author: Harshini U
Sr DBA**

switch over is executed to use the new RDS database as a primary database.

Author: Harshini U
Sr DBA

Task configuration

Task identifier
mysql-rds-migration

Descriptive Amazon Resource Name (ARN) - optional
A friendly name to override the default DMS ARN. You cannot modify it after creation.
Friendly ARN name

Replication instance
repinstance-vpc-0276c088bb1245885

Upgrades to versions 3.4.7 and higher
You have 1 instance that uses AWS DMS version 3.4.7. Upgrades to AWS DMS versions 3.4.7 and higher require that you configure AWS DMS to use VPC endpoints or use public routes. This requirement applies to source and target endpoints for these data stores: S3, Kinesis, Secrets Manager, DynamoDB, Amazon Redshift, and OpenSearch Service. [Learn more](#)

View endpoints

Source database endpoint
mysql-on-prem

Target database endpoint
rds:us-east-1

Migration type info
Migrate existing data and replicate ongoing changes

When switching database engines, the AWS Schema Conversion Tool can automatically convert your database schema and code to the engine of your choice. Click here to find out more. [Learn more](#)

Your source database is MySQL. Replicating ongoing changes requires the MySQL

- In the **Table mappings** section, choose the tables to copy over. Enter the name of the schemas and tables to copy. Use % as a wildcard character to copy multiple tables or schemas.

Table mappings

Editing mode info
Wizard
You can enter only a subset of the available table mappings.
JSON editor
You can enter all available table mappings directly in JSON format.

Specify at least one selection rule with an include action. After you do this, you can add one or more transformation rules.

Selection rules
Choose the schema and/or tables you want to include with, or exclude from, your migration task. [info](#)

Add new selection rule

where schema name is like 'test' and Source table name is like '%', include

Schema
Enter a schema

Source name
Use the % character as a wildcard
test

Source table name
Use the % character as a wildcard
%

Action
Choose "Include" to migrate your selected objects, or "Exclude" to ignore them during the migration.
Include

Source filters info
Add column filter

Transformation rules

**Author: Harshini U
Sr DBA**

4. Choose **Create task** to start your replication task.

us-east-1.console.aws.amazon.com/dms/v2/home?region=us-east-1#CreateNewTask

AWS DMS

Source name: test

Source table name: %

Action: Include

Source filters: info Add column filter

Transformation rules

Premigration assessment info

Enable premigration assessment run

Migration task startup configuration

Start migration task

☒ Automatically on create

☐ Manually later

Tags

Cancel Create task

5. After creating task, task is shown in the **Database migration tasks** section with a **Status** of **Creating**, **Starting** and then **Running**.

mysql-rds-migration created successfully.

mysql-rds-migration starting in progress.

Database migration tasks (1)


Identifier	Status	Progress	Type	Source	Target	Replication instance	Started	Stopped	Elapsed load time	Tables loaded
mysql-rds-migration	Starting	0%	Full load, ongoing replication	mysql-on-prem	firstdb	repinstance	August 1, 2022 at 12:48:24 (UTC+05:30)	-	0 s	0

Database migration tasks (1)

Identifier	Status	Progress	Type	Source	Target	Replication instance	Started	Stopped	Elapsed load time	Tables loaded
mysql-rds-migration	Running	0%	Full load, ongoing replication	mysql-on-prem	firstdb	repinstance	August 1, 2022 at 12:48:24 (UTC+05:30)	-	0 s	0

6. After the migration of existing data is completed, it shows **Status** as **Load complete, replication ongoing**. Any updates to source database at this point will be copied to target database.

Author: Harshini U
Sr DBA



7. Sample data inserted to source database after full load completed:

```
mysql> insert into tnew values (2, 'kat');
Query OK, 1 row affected (0.00 sec)

mysql> insert into tnew values (2, 'kat');
Query OK, 1 row affected (0.01 sec)

mysql> insert into tnew values (2, 'kat');
Query OK, 1 row affected (0.00 sec)

mysql> insert into tnew values (2, 'kat');
Query OK, 1 row affected (0.00 sec)

mysql> insert into tnew values (2, 'kat');
Query OK, 1 row affected (0.00 sec)

mysql> insert into tnew values (2, 'kat');
Query OK, 1 row affected (0.01 sec)

mysql> insert into tnew values (2, 'kat');
Query OK, 1 row affected (0.01 sec)

mysql> insert into tnew values (2, 'kat');
Query OK, 1 row affected (0.01 sec)

mysql> insert into tnew values (2, 'kat');
Query OK, 1 row affected (0.00 sec)

mysql> select * from tnew;
+-----+-----+
| id    | name  |
+-----+-----+
| 1     | tim   |
| 2     | kat   |
| 2     | kat   |
| 2     | kat   |
| 2     | kat   |
| 2     | kat   |
| 2     | kat   |
| 2     | kat   |
| 2     | kat   |
| 2     | kat   |
| 2     | kat   |
| 2     | kat   |
+-----+-----+
12 rows in set (0.00 sec)
```

8. Same insert data are also applied to RDS instance as well like below:

The screenshot shows the AWS DMS console for a migration task named 'mysql-rds-migration'. The 'Table statistics' tab is selected, displaying a table with migration statistics. The table has columns for Schema name, Table, Load state, Elapsed load time, Inserts, Deletes, Updates, DDLs, Applied inserts, Applied deletes, Applied updates, Applied DDLs, and Full load rows. Two rows are shown: 'test.tnew' and 'test.t1'. The 'test.tnew' row shows 10 Inserts, 0 Deletes, 0 Updates, 0 DDLs, 10 Applied inserts, 0 Applied deletes, 0 Applied updates, 0 Applied DDLs, and 2 Full load rows. The 'test.t1' row shows 0 Inserts, 0 Deletes, 0 Updates, 0 DDLs, 0 Applied inserts, 0 Applied deletes, 0 Applied updates, 0 Applied DDLs, and 0 Full load rows. The 'Inserts' and 'Full load rows' columns for the 'test.tnew' row are highlighted with red boxes.

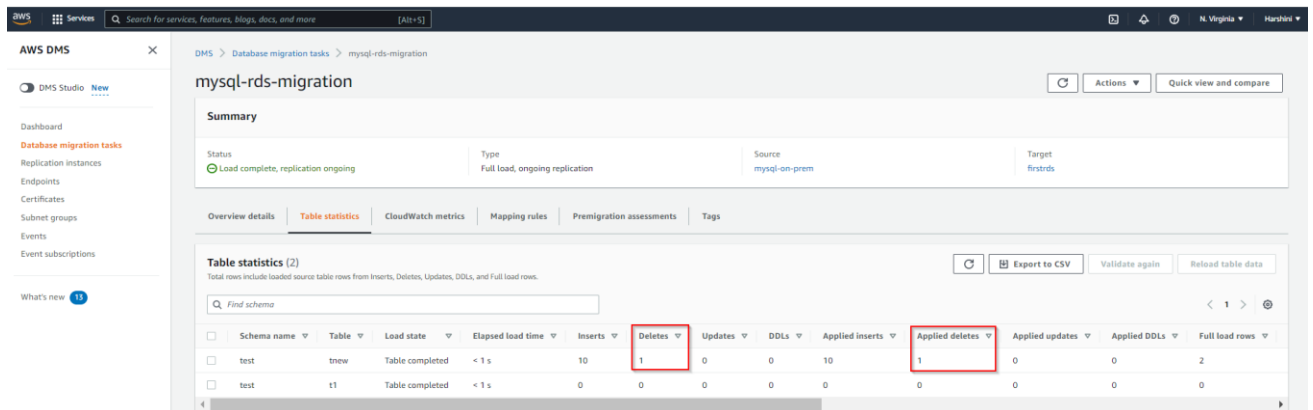
Schema name	Table	Load state	Elapsed load time	Inserts	Deletes	Updates	DDLs	Applied inserts	Applied deletes	Applied updates	Applied DDLs	Full load rows
test	tnew	Table completed	< 1 s	10	0	0	0	10	0	0	0	2
test	t1	Table completed	< 1 s	0	0	0	0	0	0	0	0	0

Author: Harshini U
Sr DBA

9. Deleted a row in the table in source database:

```
mysql> delete from tnew where id=1;
Query OK, 1 row affected (0.01 sec)
```

10. Same is reflected in the migration instance too:



The screenshot shows the AWS DMS console for a migration task named 'mysql-rds-migration'. The 'Table statistics' tab is selected, showing a table named 'tnew' with 10 inserts and 1 delete. The 'Applied deletes' column shows 1, indicating the delete operation has been replicated to the target RDS instance.

Schema name	Table	Load state	Elapsed load time	Inserts	Deletes	Updates	DDLs	Applied inserts	Applied deletes	Applied updates	Applied DDLs	Full load rows
test	tnew	Table completed	< 1 s	10	1	0	0	10	1	0	0	2
test	t1	Table completed	< 1 s	0	0	0	0	0	0	0	0	0

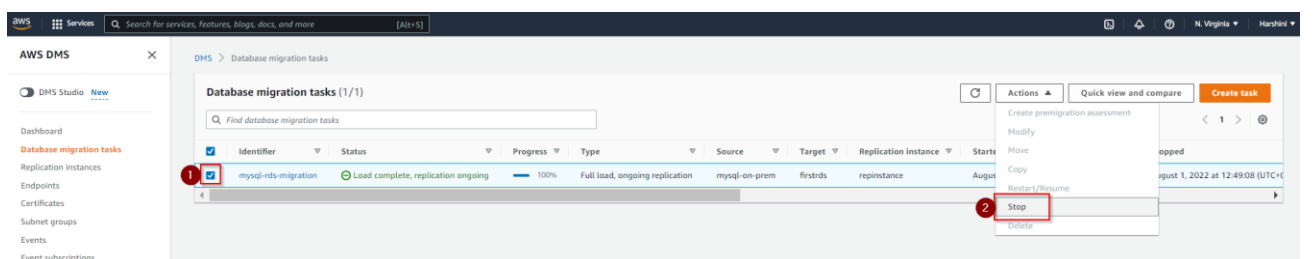
In this module, created a replication task in AWS DMS to migrate an existing data and sync ongoing changes from on-prem MySQL database to new database in Amazon RDS.

Module 5 – Clean up resources after completing migration

Now that initial migration and Change Data Capture are completed, data sync is validated in the new RDS database, RDS database is ready to use as a primary database. Also, executed a thorough test in the last module above.

After switching to use RDS database as a primary database successfully, delete AWS DMS infrastructure.

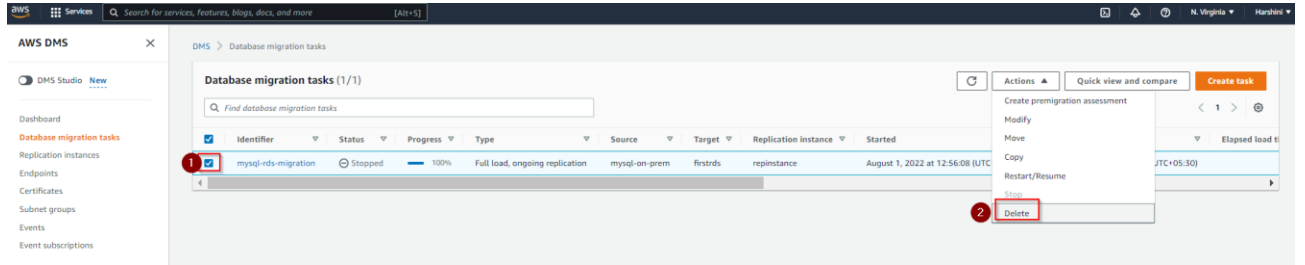
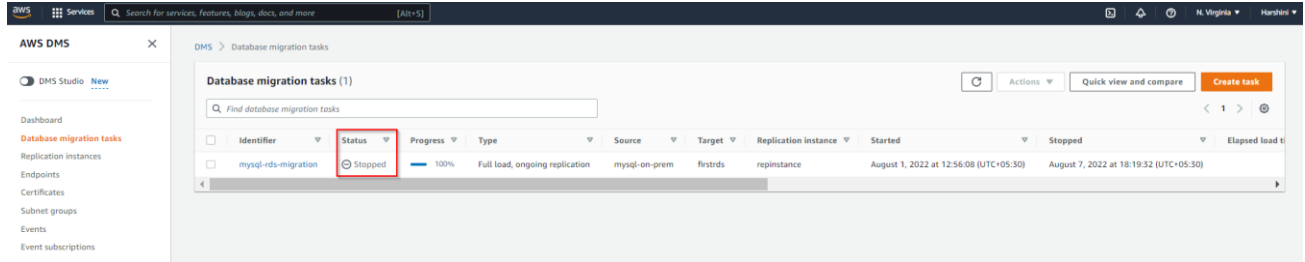
1. Stop and delete migration task like below:



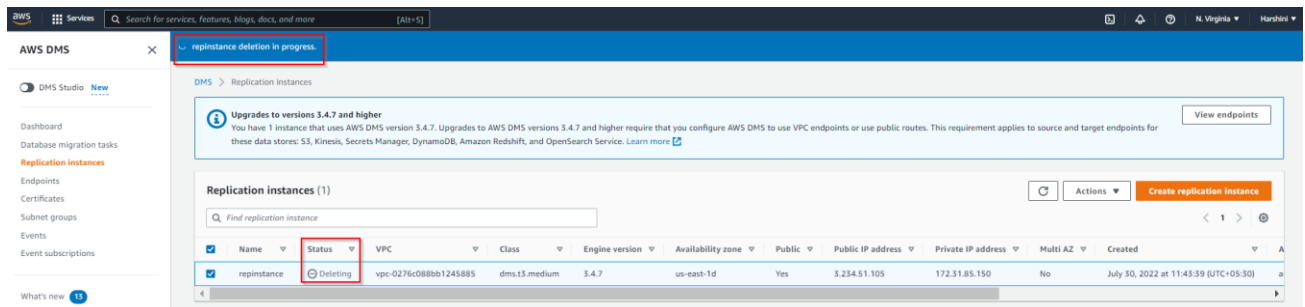
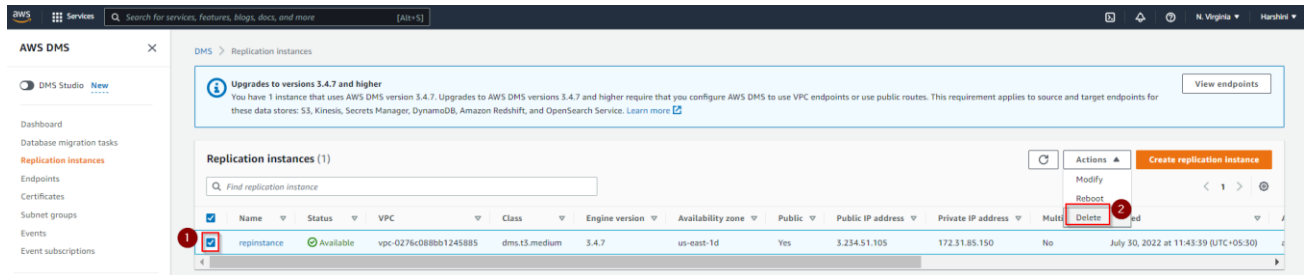
The screenshot shows the AWS DMS console for the 'mysql-rds-migration' task. The task status is 'Load complete, replication ongoing' with 100% progress. The 'Stop' button is highlighted in the 'Actions' menu.

Identifier	Status	Progress	Type	Source	Target	Replication instance	Start
mysql-rds-migration	Load complete, replication ongoing	100%	Full load, ongoing replication	mysql-on-prem	firstds	repinstance	August 1, 2022 at 12:49:08 (UTC-4)

Author: Harshini U
Sr DBA

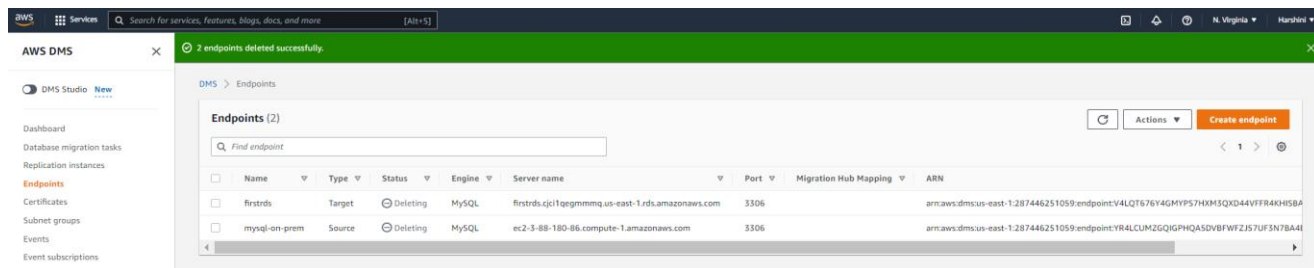
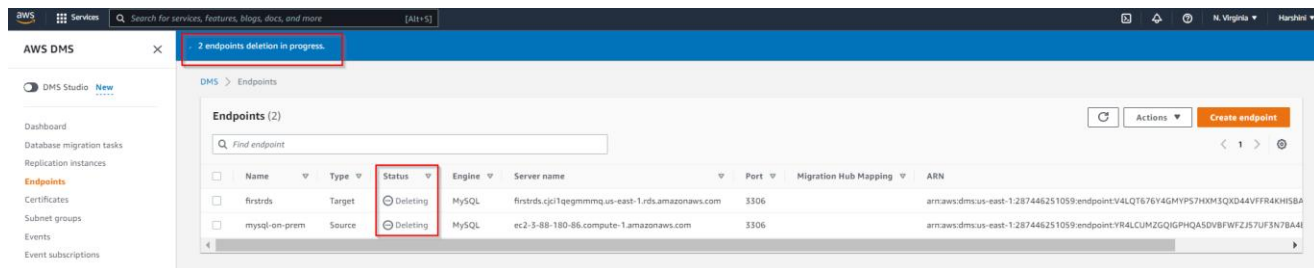

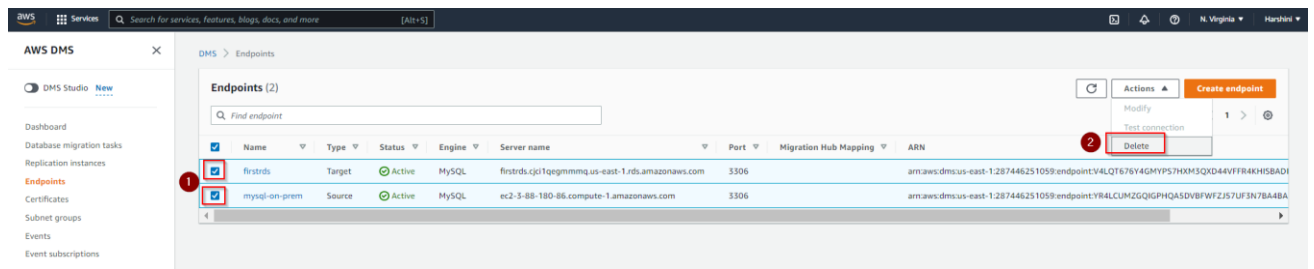


2. Select and delete replication instance:



3. Delete both the endpoints created under AWS DMS:

**Author: Harshini U
Sr DBA**



Cost Analysis

Region	Service	Upfront	Monthly	First 12 months total	Currency	Configuration Summary
US East (Ohio)	Amazon RDS for MySQL	0.00	267.00	3204.00	USD	Instance type: db.m4.xlarge Single AZ. Storage: General Purpose SSD (gp2) 100 GB

**Author: Harshini U
Sr DBA**

US East (Ohio)	Amazon Database Migration Service	0.00	64.79	777.48	USD	Instance type: dms.t3.medium (vCPU 2, Memory 4 GB, Storage: EBS only)
Total			331.79	3981.48	USD	

**Author: Harshini U
Sr DBA**

Lessons and Observations:

1. Learnt how to create Amazon RDS database instance for MySQL database.
2. Learnt to create replication instance and setup source and target endpoints in Amazon DMS.
3. Learnt how to migrate database (initial data load and change data capture) from on-prem MySQL database to RDS database instance and cutover the process after sync validation and application switchover to new RDS primary database.

**Author: Harshini U
Sr DBA**